

Digging a rabbit hole with a fork()

CocoaHeads Paris, January 2024

Frank Lefebvre

About me

- Staff Engineer @  **mojo**
- Freelance macOS, iOS & back-end developer
- @franklefebvre@mastodon.social

Context

- macOS command-line tool
- Control long-lived processes
- Single executable file

Option #1: Launch Daemon

- Apple-sanctioned way to support a long-lived process
- `/Library/LaunchDaemons/config_file.plist`
- Run at startup or on demand (`launchctl`)
- Multiple files to install
- One plist file == one process

There must be an easier way...

Option #2: UNIX fork()

- `pid_t fork(void);`
- Usually followed by `exec*()` in child branch
- `posix_spawn()`, `Process.run()`
- Parent-child relationship

Using `fork()` without `exec()`

- Parent can exit
- Child's `ppid == 1`
- Daemon

POC

- daemonize.c

```
#include <stdlib.h>
#include <unistd.h>
#include "daemonize.h"

void detachCurrentProcess(void) {
    if (fork()) { // not 0, this is the parent process
        exit(0);
    }
}
```

- Bridging-Header.h

```
#import "daemonize.h"
```


POC

CLI tool (1/3)

- swift-argument-parser package

```
import Foundation
import ArgumentParser

@main
struct UDPRecorder: ParsableCommand {
    static var configuration = CommandConfiguration(
        commandName: "udp-recorder",
        subcommands: [Create.self, StartDaemon.self]
    )

    struct Create: ParsableCommand { /* TBD */ }

    struct StartDaemon: ParsableCommand { /* TBD */ }
}
```

POC

CLI tool (2/3)

```
struct Create: ParsableCommand {
    static var configuration = CommandConfiguration(
        commandName: "create"
    )

    func run() {
        guard let path = CommandLine.arguments.first else { fatalError() }
        let url = URL(filePath: path)
        let id = UUID().uuidString
        try? Process.run(url, arguments: ["start-daemon", id]).waitUntilExit()
        print("[CREATE] done \(id)")
    }
}
```

POC

CLI tool (3/3)

```
struct StartDaemon: ParsableCommand {
    static var configuration = CommandConfiguration(
        commandName: "start-daemon",
        shouldDisplay: false
    )
    @Argument var id: String

    func run() {
        print("[DAEMON] enter - pid = \(ProcessInfo.processInfo.processIdentifier) - id = \(id)")
        detachCurrentProcess()
        print("[DAEMON] child - pid = \(ProcessInfo.processInfo.processIdentifier) - id = \(id)")
        Thread.sleep(forTimeInterval: 20)
        print("[DAEMON] end")
    }
}
```

POC

Testing...



POC

Testing...



It works!



Integration

Communication layer: UNIX local socket

```
import Foundation
import KituraNet

final class UnixSocketListener {
    private var server: HTTPServer? = nil

    init(id: String) {
        do {
            try FileManager.default.createDirectory(at: UnixSocket.url,
withIntermediateDirectories: true)
            self.server = try HTTPServer.listen(unixDomainSocketPath: "\(UnixSocket.path)/\(id)",
delegate: self)
        } catch {
            print("UnixSocketListener error \(error.localizedDescription)")
        }
    }
}

extension UnixSocketListener: ServerDelegate { /* TBD */ }
```

Integration

Using ArgumentParser with async/await

```
@main
struct UDPRecorder: ParsableCommand {
    static var configuration = CommandConfiguration(
        commandName: "udp-recorder",
        subcommands: [Create.self, StartDaemon.self]
    )

    struct Create: ParsableCommand {
        static var configuration = CommandConfiguration(
            commandName: "create"

            func run() {
                /* implementation */
            }
        }

        /* struct StartDaemon */
    }
}
```


Integration

Using ArgumentParser with async/await

```
@main
struct UDPRecorder: ParsableCommand {
    static var configuration = CommandConfiguration(
        commandName: "udp-recorder",
        subcommands: [Create.self, StartDaemon.self]
    )

    struct Create: ParsableCommand {
        static var configuration = CommandConfiguration(
            commandName: "create")

        func run() async {
            /* implementation */
        }
    }

    /* struct StartDaemon */
}
```

Integration

Using ArgumentParser with async/await

```
@main
struct UDPRecorder: AsyncParsableCommand {
    static var configuration = CommandConfiguration(
        commandName: "udp-recorder",
        subcommands: [Create.self, StartDaemon.self]
    )

    struct Create: AsyncParsableCommand {
        static var configuration = CommandConfiguration(
            commandName: "create")

        func run() async {
            /* implementation */
        }
    }

    /* struct StartDaemon */
}
```

Testing...

Testing...

```
/tmp/udp-recorder create
```

Testing...

```
/tmp/udp-recorder create
```

```
+ [Swift.__SharedStringStorage initialize] may have been in progress in another thread  
when fork() was called. We cannot safely call it or ignore it in the fork() child  
process. Crashing instead. Set a breakpoint on objc_initializeAfterForkError to debug.
```

Testing...

```
/tmp/udp-recorder create
```

```
+ [Swift.__SharedStringStorage initialize] may have been in progress in another thread  
when fork() was called. We cannot safely call it or ignore it in the fork() child  
process. Crashing instead. Set a breakpoint on objc_initializeAfterForkError to debug.
```

```
OBJC_DISABLE_INITIALIZE_FORK_SAFETY=YES /tmp/udp-recorder create
```

Testing...

```
/tmp/udp-recorder create
```

```
+ [Swift.__SharedStringStorage initialize] may have been in progress in another thread  
when fork() was called. We cannot safely call it or ignore it in the fork() child  
process. Crashing instead. Set a breakpoint on objc_initializeAfterForkError to debug.
```

```
OBJC_DISABLE_INITIALIZE_FORK_SAFETY=YES /tmp/udp-recorder create
```

Can't communicate with the child process...

Testing...

```
/tmp/udp-recorder create
```

```
+ [Swift.__SharedStringStorage initialize] may have been in progress in another thread  
when fork() was called. We cannot safely call it or ignore it in the fork() child  
process. Crashing instead. Set a breakpoint on objc_initializeAfterForkError to debug.
```

```
OBJC_DISABLE_INITIALIZE_FORK_SAFETY=YES /tmp/udp-recorder create
```

Can't communicate with the child process...

... because the child has crashed

Testing...

```
/tmp/udp-recorder create
```

```
+ [Swift.__SharedStringStorage initialize] may have been in progress in another thread  
when fork() was called. We cannot safely call it or ignore it in the fork() child  
process. Crashing instead. Set a breakpoint on objc_initializeAfterForkError to debug.
```

```
OBJC_DISABLE_INITIALIZE_FORK_SAFETY=YES /tmp/udp-recorder create
```

Can't communicate with the child process...

- ... because the child has crashed

- ... in libDispatch when starting the server

Testing...

```
/tmp/udp-recorder create
```

```
+ [Swift.__SharedStringStorage initialize] may have been in progress in another thread  
when fork() was called. We cannot safely call it or ignore it in the fork() child  
process. Crashing instead. Set a breakpoint on objc_initializeAfterForkError to debug.
```

```
OBJC_DISABLE_INITIALIZE_FORK_SAFETY=YES /tmp/udp-recorder create
```

Can't communicate with the child process...

... because the child has crashed

... in libDispatch when starting the server

The process has forked and you cannot use this CoreFoundation functionality safely.
You MUST exec(). Break on
__THE_PROCESS_HAS_FORKED_AND_YOU_CANNOT_USE_THIS_COREFOUNDATION_FUNCTIONALITY___YOU_MU
ST_EXEC__() to debug.

Why?

<https://forums.developer.apple.com/forums/thread/737464>

Combining `fork` with Apple's frameworks is a tricky business [1]. If you only use Posix APIs, `fork` should behave reliably. Once you start using our higher-level frameworks, well... we try to keep things working but we can't make any guarantees.

There are two cases here:

- Calling `fork` without `exec*` to create a long-running 'clone' of the current process
- Combining `fork` and `exec*` to run a new program in a child process

The first case will (usually :-)) work if you limit yourself to Posix APIs. It's fundamentally incompatible with our higher-level frameworks.

For the second case, the best way to avoid problems is to use an API that combines `fork` and `exec*`. At the BSD level, that's `posix_spawn`. Higher up, you have `NSTask` (aka `Process` in Swift).




Share and Enjoy

—

Quinn "The Eskimo!" @ Developer Technical Support @ Apple

```
let myEmail = "eskimo" + "1" + "@" + "apple.com"
```

[1] There's a fundamental disconnect between BSD and Mach on this topic, and Apple's frameworks rely on Mach a *lot*.

Posted 2 months ago by   eskimo 

Why?

<https://forums.developer.apple.com/forums/thread/737464>

Combining `fork` with Apple's frameworks is a tricky business [1]. If you only use Posix APIs, `fork` should behave reliably. Once you start using our higher-level frameworks, well... we try to keep things working but we can't make any guarantees.

There are two cases here:

- Calling `fork` without `exec*` to create a long-running 'clone' of the current process
- Combining `fork` and `exec*` to run a new program in a child process




The first case will (usually :-)) work if you limit yourself to Posix APIs. It's fundamentally incompatible with our higher-level frameworks.

For the second case, the best way to avoid problems is to use an API that combines `fork` and `exec*`. At the BSD level, that's `posix_spawn`. Higher up, you have `NSTask` (aka `Process` in Swift).

Share and Enjoy

Quinn "The Eskimo!" @ Developer Technical Support @ Apple
`let myEmail = "eskimo" + "1" + "@" + "apple.com"`

[1] There's a fundamental disconnect between BSD and Mach on this topic, and Apple's frameworks rely on Mach a *lot*.

Posted 2 months ago by   eskimo 

Why?

<https://forums.developer.apple.com/forums/thread/737464>

Combining `fork` with Apple's frameworks is a tricky business [1]. If you only use Posix APIs, `fork` should behave reliably. Once you start using our higher-level frameworks, well... we try to keep things working but we can't make any guarantees.

There are two cases here:

- Calling `fork` without `exec*` to create a long-running 'clone' of the current process
- Combining `fork` and `exec*` to run a new program in a child process

The first case will (usually :-)) work if you limit yourself to Posix APIs. It's fundamentally incompatible with our higher-level frameworks.

For the second case, the best way to avoid problems is to use an API that combines `fork` and `exec*`. At the BSD level, that's `posix_spawn`. Higher up, you have `NSTask` (aka `Process` in Swift).




Share and Enjoy

—

Quinn "The Eskimo!" @ Developer Technical Support @ Apple

```
let myEmail = "eskimo" + "1" + "@" + "apple.com"
```

[1] There's a fundamental disconnect between BSD and Mach on this topic, and Apple's frameworks rely on Mach a *lot*.

Posted 2 months ago by   eskimo 

Why?

<https://forums.developer.apple.com/forums/thread/737464>

Combining `fork` with Apple's frameworks is a tricky business [1]. If you only use Posix APIs, `fork` should behave reliably. Once you start using our higher-level frameworks, well... we try to keep things working but we can't make any guarantees.

There are two cases here:

- Calling `fork` without `exec*` to create a long-running 'clone' of the current process
- Combining `fork` and `exec*` to run a new program in a child process

The first case will (usually :-) work if you limit yourself to Posix APIs. It's fundamentally incompatible with our higher-level frameworks.

For the second case, the best way to avoid problems is to use an API that combines `fork` and `exec*`. At the BSD level, that's `posix_spawn`. Higher up, you have `NSTask` (aka `Process` in Swift).


Share and Enjoy

—

Quinn "The Eskimo!" @ Developer Technical Support @ Apple

```
let myEmail = "eskimo" + "1" + "@" + "apple.com"
```

[1] There's a fundamental disconnect between BSD and Mach on this topic, and Apple's frameworks rely on Mach a *lot*.

Posted 2 months ago by   eskimo 

Option #3

Trampoline CLI tool

- POSIX-only
- daemonize + `posix_spawn` multithreaded CLI tool

Option #3

Trampoline CLI tool

- POSIX-only
- daemonize + `posix_spawn` multithreaded CLI tool
- This tool exists: `/usr/bin/nohup`

Final Code

```
struct Create: AsyncParsableCommand {
    static var configuration = CommandConfiguration(commandName: "create")

    func run() async {
        let id = UUID().uuidString
        guard let path = CommandLine.arguments.first else { fatalError() }
        let command = URL(filePath: "/usr/bin/nohup")
        let arguments = [path, "start-recorder", id]
        do {
            let process = Process()
            process.executableURL = command
            process.arguments = arguments
            process.standardOutput = nil
            process.standardError = nil
            try process.run()
        } catch {
            print("Process failed: \(error)")
        }
        // ... send request to local socket and wait for response
    }
}
```

Questions?